

# Statistical Learning

Due in two weeks on Moodle (May 31 + tolerance)

## Homework-02

### General Instructions

- You can use *any* programming language you want, as long as your work is runnable/correct/readable. Two examples:
  - **In R:** it would be nice to upload a well-edited and working **R Markdown** file (`.rmd`) + its **html** output.
  - **In Python:** it would be nice to upload a well-edited and working **Jupyter notebook** (or similia).
- Remember our **policy on collaboration**:

*Collaboration on homework assignments with fellow students is **encouraged**.*

*However, such collaboration should be clearly acknowledged, by listing the names of the students with whom you have had discussions concerning your solution.*

*You may **not**, however, share written work or code after discussing a problem with others.*

*The solutions should be written by **you**.*

### In case of R

If you go for R, to be sure that everything is working, start **RStudio** and create an empty project called **HW1**. Now open a new **R Markdown** file (**File > New File > R Markdown...**); set the output to **HTML mode**, press **OK** and then click on **Knit HTML**. This should produce a **html**. You can now start editing this file to produce your homework submission.

- For more info on **R Markdown**, check the support webpage: [R Markdown from RStudio](#).
- For more info on how to write math formulas in LaTeX: [Wikibooks](#).

### Exercise 1: The Bayes Classifier

#### ↪ Your job ←

Suppose that  $(Y, X)$  are random variable with  $Y \in \{0, 1\}$  and  $X \in \mathbb{R}$ . Suppose that

$$(X | Y = 0) \sim \text{Unif}(-3, 1) \quad \text{and} \quad (X | Y = 1) \sim \text{Unif}(-1, 3).$$

Further suppose that  $\mathbb{P}(Y = 0) = \mathbb{P}(Y = 1) = \frac{1}{2}$ .

1. Define the Bayes classifier/strategy and briefly explain its role/importance in classification/prediction.
2. Find (with pen and paper) the Bayes classification rule  $h_{\text{opt}}(x)$ .
3. Simulate  $n = 250$  data from the joint data model  $p(y, x) = p(x | y) \cdot p(y)$  described above, and then:
  - Plot the data together with the regression function that defines  $h_{\text{opt}}(x)$
  - Evaluate the performance of the Bayes Classifiers on these simple (only 1 feature!) data
  - Apply any other classifier of your choice to these data and comparatively comment its performance (...with respect to those of the Bayes classifiers). Of course those  $n = 250$  training data should be used for training and validation too (in case there are tuning-parameters)
4. Since you are simulating the data, you can actually see what happens in repeated sampling. Hence, repeat the sampling  $M = 1000$  times keeping  $n = 250$  fixed (a simple **for**-loop will do it), and redo the comparison. Who's the best now? Comment.

## Exercise 2: Go Gradient, Go!

For this exercise you will be working on a dataset of **product reviews on Amazon** that I've pre-processed in order to save you some time (see the box below for the details, also [linked here](#)). The dataset consists of a random sample of reviews of **books** and an equal number of reviews from **film/television series**.

More specifically there are ~150,000 training data and ~50,000 test data with 1334 features/words.

Your goal is to build models that predicts whether the reviews is of a **books** or a **film/television series** using word counts as features vectors.

The data are stored in a single *R-friendly* file called **amazon\_review\_clean.RData** and everything, response vector and feature matrix, are already splitted into training and test sets: **y\_tr**, **X\_tr**, **y\_te** and **X\_te**. Since these are very simple, *tidy* objects, it's easy to export them as ASCII files with **write.table()** to start working in **Python** or any other programming language.

The **X**-matrices are known as *term-frequency matrices* or *document-term matrices* and count how often various words are used in the text. For more info have also a look at the **Kernel for Text** slide set (pp. 142-147).

Specifically, the **X[i,j]** element counts the number of times the word  $w_j$  appears in the  $i^{\text{th}}$  review for some pre-specified set of words of interest. These matrices were produced using the **cleanNLP** package (see the box below for details): I included any word that occurred in at least 0.1% and no more than 50% of the reviews for a total of about 1300 words.

### ↪ Your job ↩

1. Looking at the formula at the bottom of **page 8 of our notes**, train a linear classifier by gradient descent (GD). Check its performance and then comment the results.

HINT: Please notice that you can pre-compute (meaning, compute *before* the GD loop) the heavy part.

2. **Bonus.** If you are brave enough, go **stochastic** (page 40-43) and comment your "*experience*"...

---

## Additional Code / Data Cleaning and Preparation

```
# Packages -----
require(readr)
require(cleanNLP) # https://statmaths.github.io/cleanNLP/
require(stringi)
require(caret)

# Load & Look -----
amazon <- read_csv("amazon.csv")
str(amazon)
table(amazon$class)
stri_wrap(amazon$text[amazon$class == "book"][1:10])

# Remove NA otherwise troubles with <cnlp_annotates> below
len <- stri_length(amazon$text)
hist(len, breaks = 1000)
summary(len) # wow! min length = 2, average length = 917, max length...30k!! XD
quantile(len, seq(0,1,.1), na.rm = T)

idx_na <- which(is.na(len))
amazon <- amazon[-idx_na,]

# Clean -----
?cnlp_init_stringi
?cnlp_annotate

cnlp_init_stringi() # initialize tokenizer backend
anno <- cnlp_annotate(amazon, text_name = "text") # take some time...

# tf-idf score -----
```

```

# See our notes: Kernel for Text (pp. 142-147)
# https://elearning.uniroma1.it/pluginfile.php/1029332/mod_folder/intro/Lecture_11.pdf
?cnlp_utils_tfidf

# The options in the call determine what words are included:
# a word must be used in at least <min_df> percent of documents
# but not in more than <max_df> documents
X <- cnlp_utils_tfidf(anno$token,
                      min_df = 0.01, max_df = 0.5,
                      tf_weight = "raw")

# Take a look
dim(X)
colnames(X)
round(as.matrix(X[1:10, 1:10]), 2)

# Train-Test split -----

y <- amazon$class
X <- as.matrix(X)

set.seed(124) # for reproducibility
idx_tr <- createDataPartition(y = y, p = .75, list = F)

X_tr <- X[ idx_tr, ]; X_te <- X[-idx_tr, ]
y_tr <- y[ idx_tr ]; y_te <- y[-idx_tr]

dim(X_tr); dim(X_te); length(y_tr); length(y_te)

```